



1. A good problem to solve using a regression supervised learning technique would be item pricing in retail based on the current time of day/month/year (for example, Christmas items are more popular around November so their prices can be increased to generate more revenue).
  - a. As features (x), I would use the time of day/month/year.
  - b. As labels (y), I would use the price of items.
  - c. I would collect data by conducting a trial over some period given existing price trends, and I would also consider artificially fluctuating prices to see how the market responds.
  - d. This problem can be challenging because there are multiple factors to consider that can affect a buyer's shopping habits, and it will be unclear if the current time alone is enough to affect habits.
2. A good problem to solve using a classification supervised learning technique would be speech recognition to classify certain sounds as words.
  - a. As features (x), I would use the sound input.
  - b. As labels (y), I would use the words that the sounds represent.
  - c. I would collect data by providing various samples of recordings along with the words that they represent.
  - d. This problem can be difficult because there are infinite ways that a word can be represented in sound, due to accents, tones of voice, and other factors.

3c.

The histogram for x does look like a Gaussian distribution.



The histogram for z does look like a uniform distribution.



3d.

```
Loop execution time: 1.1336390972137451
```

3e.

```
Numpy function time: 0.00015497207641601562
```

The NumPy function is far more efficient.

3f.

Trial 1:

```
Number of elements retrieved: 374828
```

Trial 2:

```
Number of elements retrieved: 374627
```

Trial 3:

```
Number of elements retrieved: 374426
```

There is a slight difference, but the numbers always float around 375000. This is because we used a randomization function to choose the numbers, but the given range will on average account for about 37.5% of the data because we have used a uniform distribution.

4a.

```
4a results:
Minimum value in each column:
[[2 1 3]]
Max value in each row:
[[ 3]
 [ 8]
 [18]]
Highest value in A:
18
Sum of columns:
[[10 15 29]]
Sum of A:
54
B:
[[ 4  1  9]
 [ 4 36 64]
 [36 64 324]]
```

4b.

```
Solution for 4b:
[[0.3]
 [0.4]
 [0. ]]
```

4c.

4c.)

$$x_1 = [-0.5 \ 0 \ 1.5]$$

$$|x_1|_1 = 0.5 + 0 + 1.5 = 2$$

$$|x_1| = \sqrt{(0.5)^2 + (0)^2 + (1.5)^2} = 1.58$$

$$x_2 = [-1 \ -1 \ 0]$$

$$|x_2|_1 = 1 + 1 + 0 = 2$$

$$|x_2| = \sqrt{(1)^2 + (1)^2 + (0)^2} = 1.41$$

```

4c NumPy calculations:
x1_l1: 2.0
x1_l2: 1.5811388300841898
x2_l1: 2.0
x2_l2: 1.4142135623730951

```

5a.

```

5a X:
[[ 1  1  1]
 [ 2  2  2]
 [ 3  3  3]
 [ 4  4  4]
 [ 5  5  5]
 [ 6  6  6]
 [ 7  7  7]
 [ 8  8  8]
 [ 9  9  9]
 [10 10 10]]

```

Trial 1:	Trial 2:	Trial 3:
<pre> X_train: [[ 7  7  7]  [ 8  8  8]  [ 6  6  6]  [ 1  1  1]  [ 4  4  4]  [10 10 10]  [ 2  2  2]  [ 5  5  5]] X_test: [[9 9 9]  [3 3 3]] y_train: [[ 7]  [ 8]  [ 6]  [ 1]  [ 4]  [10]  [ 2]  [ 5]] y_test: [[9]  [3]] </pre>	<pre> X_train: [[ 1  1  1]  [10 10 10]  [ 6  6  6]  [ 5  5  5]  [ 8  8  8]  [ 2  2  2]  [ 4  4  4]  [ 3  3  3]] X_test: [[9 9 9]  [7 7 7]] y_train: [[ 1]  [10]  [ 6]  [ 5]  [ 8]  [ 2]  [ 4]  [ 3]] y_test: [[9]  [7]] </pre>	<pre> X_train: [[ 2  2  2]  [ 4  4  4]  [ 8  8  8]  [ 5  5  5]  [ 9  9  9]  [ 3  3  3]  [10 10 10]  [ 1  1  1]] X_test: [[6 6 6]  [7 7 7]] y_train: [[ 2]  [ 4]  [ 8]  [ 5]  [ 9]  [ 3]  [10]  [ 1]] y_test: [[6]  [7]] </pre>

We do not get the same submatrices each time. This is because, again, we used a randomization function to shuffle the matrix rows to choose the submatrices. Thus, we should expect different results each time.